



## Computer Vision Challenge Track: Target and Eliminate

### Problem:

Weeds to increase crop yields Weeds are an unwanted intruder in the agricultural business. They steal nutrients, water, land, and other critical resources to grow healthy crops. These intruders can lead to lower yields and inefficient deployment of resources by farmers. One known approach is to use pesticides to remove weeds, but aggressive pesticides create health risks for humans. Computer vision technology can automatically detect the presence of weeds and use targeted remediation techniques to remove them from fields with minimal environmental impact.

### Expected Solution:

In this hackathon track, you will be tasked with training and deploying a model into a simulated production environment - where your binary-classification accuracy (F1 score) and inference time will be used to rank you against other teams competing for this track's top spot.

---

### We leveraged the following Intel® AI Analytics Toolkit (AI Kit) - OneAPI Libraries for this model development

1. Intel® Optimization for TensorFlow\*
2. Intel® Neural Compressor\*
3. Intel® Optimization for PyTorch\*

---

**This has greatly reduced the time of our overall processing compared to standard libraries**

---

## Importing IntelTensorflow API

**intel/intel-extension-for-tensorflow**

Intel® Extension for TensorFlow\*



```
In [7]: 1 #pip install --upgrade intel-extension-for-tensorflow[cpu]
2 import tensorflow as tf
3
4 from neural_compressor.config import PostTrainingQuantConfig
5 from neural_compressor.data import DataLoader
6 from neural_compressor.data import Datasets
7
8 print(tf.__version__)
```

2.11.0

---

## Importing all the required Python libraries (Intel API and others)

```
In [8]: 1 import pandas as pd
2 import matplotlib.pyplot as plt
3 import seaborn as sns
4 import os
5 import shutil
6 import glob
7 import cv2
8 import numpy as np
9 import math
10 import datetime
11 import time
12 import random
13 import gc
14 from tqdm import tqdm
15 from IPython.display import Markdown, display
```

```
In [9]: 1 import keras
2 import tensorflow as tf
3 from sklearn import metrics
4 from keras.models import Sequential
5 from tensorflow.keras.layers import Dense, Dropout, Activation, Flatten, Conv2D, Max
6 from keras.preprocessing.image import ImageDataGenerator
7 from keras.optimizers import Adam
8 from sklearn.metrics import classification_report, confusion_matrix
9 from tensorflow.keras import optimizers
10 from tensorflow.keras import applications
11 from sklearn.metrics import confusion_matrix, classification_report
12 from sklearn.datasets import make_circles
13 from sklearn.metrics import accuracy_score
14 from sklearn.metrics import precision_score
15 from sklearn.metrics import recall_score
16 from sklearn.metrics import f1_score
17 from sklearn.metrics import cohen_kappa_score
18 from sklearn.metrics import roc_auc_score
19 from sklearn.model_selection import train_test_split
```

---

## Preparing the Dataset

From the original **dataset ( data folder )** extracting the images in two different folders called **'No\_Weed'** and **'Weed'** based on the labels given in the text file related to the Images. Images with **Class 0** goes to **'No\_Weed'** folder and Images with **Class 1** goes to **'Weed'** folder

```
In [10]: 1 path = 'Dataset/data/'
2 target_0 = 'Dataset/No_Weed/'
3 target_1 = 'Dataset/Weed/'
4 ctr = 0
5 if not os.path.isdir(target_0) and not os.path.isdir(target_1):
6     os.mkdir(target_0)
7     os.mkdir(target_1)
8     for i in range(1, len(os.listdir(path)), 2):
9         file = os.listdir(path)[i]
10        img = os.listdir(path)[i - 1]
11        abs_path = path + file
12        f = open(abs_path, 'r')
13        target = f.read()
14        if int(target[0]) == 0:
15            src_path = path + img
16            shutil.copy(src_path, target_0)
17        else:
18            src_path = path + img
19            shutil.copy(src_path, target_1)
```

---

## Loading the Dataset

- Dividing the dataset based on labels. `data_0` contains all the images which are not infected with weed and `data_1` contains all the images which are infected with weed and likewise for test data also.

```
In [11]: 1 from sklearn import datasets
          2 dataset = 'Dataset'
          3 #TEST_DIR = f'{DATADIR}/data/test'
          4 #IMG_SIZE = 50
          5 #LR = 1e-3
          6
          7 #Getting Train Data - No Weed and Weed
          8 data_0 = [dataset + '/No_Weed' + '/' + '{}'.format(i)
          9             for i in os.listdir(dataset + '/No_Weed')]
         10 data_1 = [dataset + '/Weed' + '/' + '{}'.format(i)
         11             for i in os.listdir(dataset + '/Weed')]
         12
         13
         14 imgs_data = data_0 + data_1
         15 #random.shuffle(train_imgs) # shuffle it randomly
```

---

## Data Pre-Processing

- function to read and process the images to an acceptable format for our machine learning model

In [12]:

```
1 nrows = 224
2 ncolumns = 224
3 #channels = 3 #change to 1 if you want to use grayscale image
4
5
6 #A function to read and process the images to an acceptable format for our model
7 def read_and_process_image(list_of_images):
8     """
9     Returns two arrays:
10     X is an array of resized images
11     y is an array of labels
12     """
13
14     X = [] # images
15     y = [] # labels
16
17     for image in tqdm(list_of_images):
18         X.append(cv2.resize(cv2.imread(image, cv2.IMREAD_COLOR),
19                             (nrows,ncolumns), interpolation=cv2.INTER_CUBIC)) #Read
20         #get the labels
21         if 'No_Weed' in image[:15]:
22             y.append(0)
23         elif 'Weed' in image[:15]:
24             y.append(1)
25
26     return X, y
```

In [13]:


```
1 X, y = read_and_process_image(imgs_data)
```

```
100%|████████████████████████████████████████████████████████████████████████████████| 1
301/1301 [00:19<00:00, 68.10it/s]
```

In [14]:

```
1 plt.figure(figsize=(20,10))
2 columns = 5
3 for i in tqdm(range(columns)):
4     plt.subplot(5 // columns + 1, columns, i + 1)
5     plt.imshow(X[i])
```

```
100%|████████████████████████████████████████████████████████████████████████████████|
| 5/5 [00:00<00:00, 35.57it/s]
```

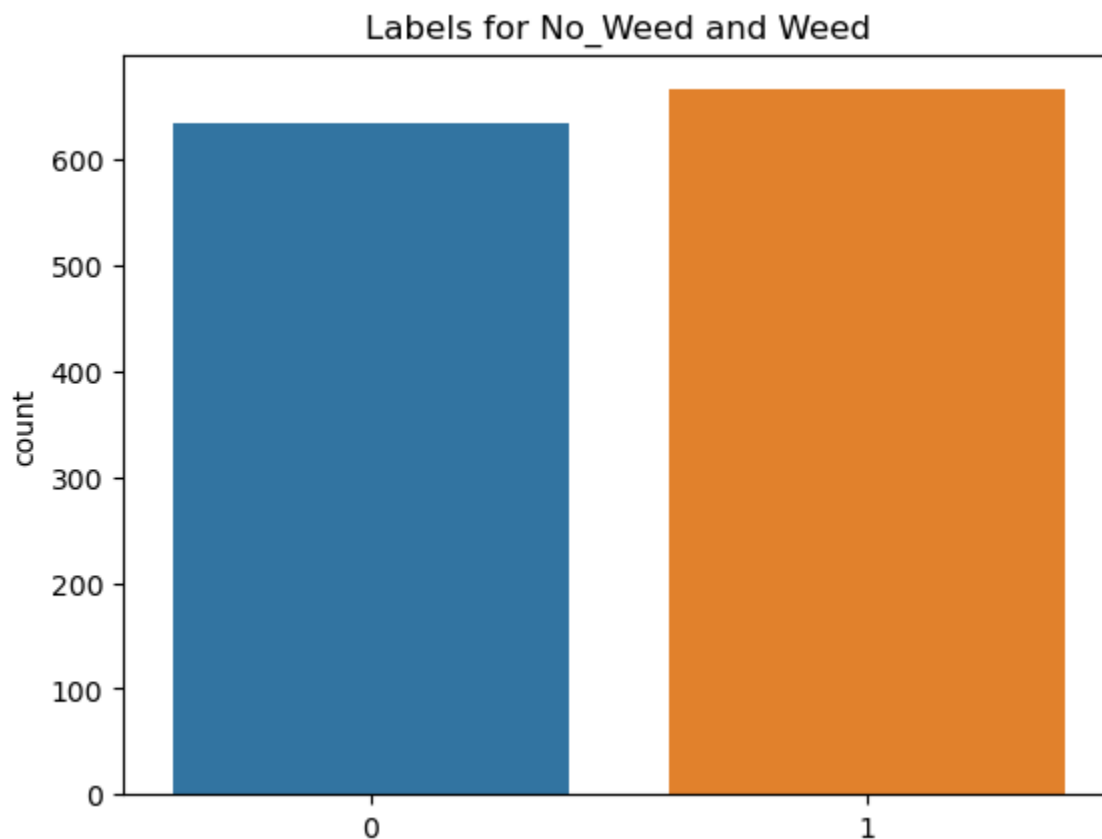


```
In [15]: 1 del imgs_data
2 gc.collect()
3
4 #Convert list to numpy array
5 X = np.array(X)
6 y = np.array(y)
7
8
9 #Lets plot the label to be sure we just have two class
10 sns.countplot(y)
11 plt.title('Labels for No_Weed and Weed')
```

C:\Users\ashutoshvmadmin\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(

Out[15]: Text(0.5, 1.0, 'Labels for No\_Weed and Weed')



---

## Splitting the Data into Training and Testing Set

```
In [16]: 1 X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.20, random_stat
2
3 print("Shape of train images is:", X_train.shape)
4 print("Shape of validation images is:", X_val.shape)
5 print("Shape of labels is:", y_train.shape)
6 print("Shape of labels is:", y_val.shape)
```

```
Shape of train images is: (1040, 224, 224, 3)
Shape of validation images is: (261, 224, 224, 3)
Shape of labels is: (1040,)
Shape of labels is: (261,)
```

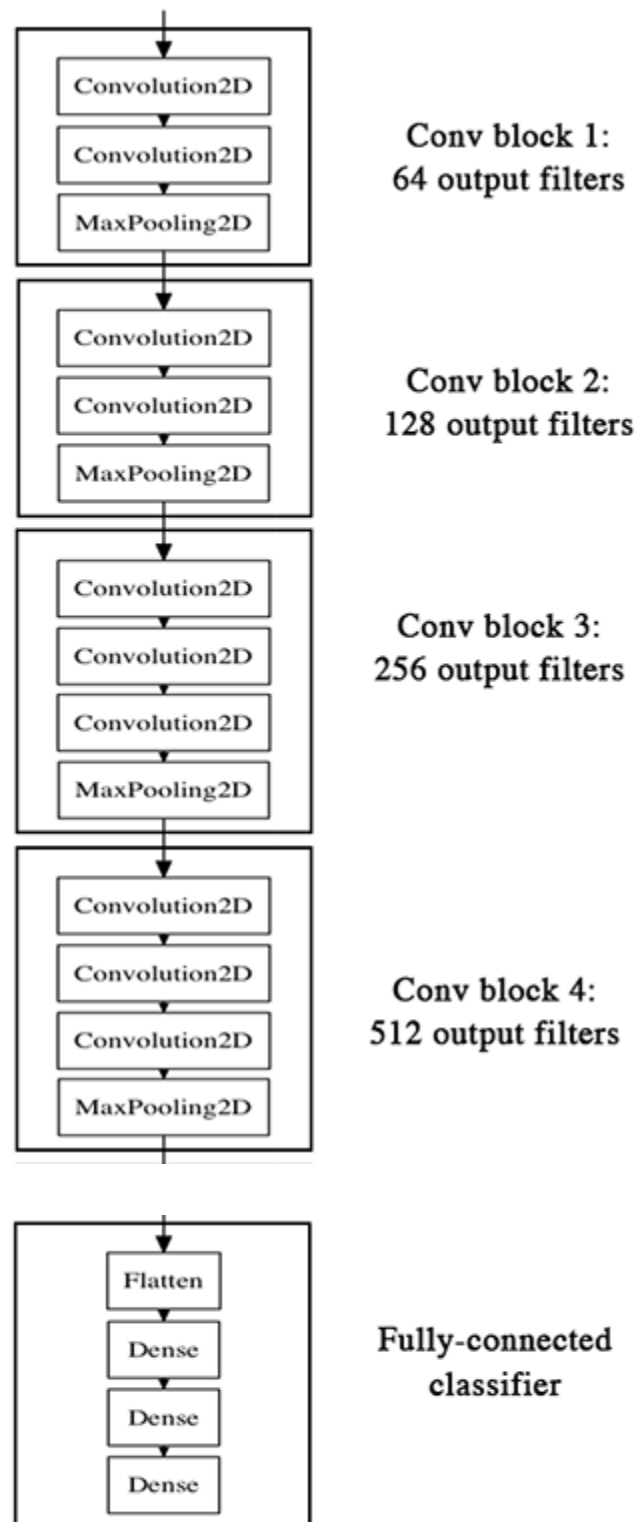
```
In [17]: 1 print("Shape of train images is:", X.shape)
2 print("Shape of labels is:", y.shape)
```

```
Shape of train images is: (1301, 224, 224, 3)
Shape of labels is: (1301,)
```

```
In [18]: 1 del X
2 del y
3 gc.collect()
4
5 #get the length of the train and validation data
6 ntrain = len(X_train)
7 nval = len(X_val)
8 print(ntrain)
9 print(nval)
10
11 #We will use a batch size of 32. Note: batch size should be a factor of 2.**4,8,16,
12 batch_size = 32
```

```
1040
261
```

## Preparing our Initial Model for Training





```
In [19]: 1 model = Sequential()
2
3 model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(224, 224, 3)))
4 model.add(MaxPooling2D((2, 2)))
5
6 model.add(Conv2D(64, (3, 3), activation='relu'))
7 model.add(MaxPooling2D((2, 2)))
8
9 model.add(Conv2D(128, (3, 3), activation='relu'))
10 model.add(MaxPooling2D((2, 2)))
11
12 model.add(Conv2D(256, (3, 3), activation='relu'))
13 model.add(MaxPooling2D((2, 2)))
14
15 model.add(Flatten())
16 model.add(Dropout(0.5)) #Dropout for regularization
17 model.add(Dense(512, activation='relu'))
18 model.add(Dense(1, activation='sigmoid'))
```

```
In [20]: 1 model.compile(loss='binary_crossentropy', optimizer=optimizers.RMSprop(lr=1e-4), me
```

WARNING:absl:`lr` is deprecated, please use `learning\_rate` instead, or use the legacy optimizer, e.g., tf.keras.optimizers.legacy.RMSprop.

```
In [21]: 1 train_datagen = ImageDataGenerator(rescale=1./255, #Scale the image between 0 and
2                                     rotation_range=50,
3                                     width_shift_range=0.2,
4                                     height_shift_range=0.2,
5                                     shear_range=0.2,
6                                     zoom_range=0.2,
7                                     horizontal_flip=True,)
8
9 #added
10 train_datagen = ImageDataGenerator(rescale=1./255 #Scale the image between 0 and
11                                     )
12
13
14 val_datagen = ImageDataGenerator(rescale=1./255) #We do not augment validation data
```

```
In [22]: 1 train_generator = train_datagen.flow(X_train, y_train, batch_size=batch_size)
2 val_generator = val_datagen.flow(X_val, y_val, batch_size=batch_size)
```



Epoch 1/50  
32/32 [=====] - 39s 1s/step - loss: 0.7865 - acc: 0.5169 - val\_loss: 0.6841 - val\_acc: 0.5391  
Epoch 2/50  
32/32 [=====] - 38s 1s/step - loss: 0.7231 - acc: 0.5893 - val\_loss: 0.5618 - val\_acc: 0.7148  
Epoch 3/50  
32/32 [=====] - 37s 1s/step - loss: 0.6851 - acc: 0.6240 - val\_loss: 0.4502 - val\_acc: 0.7969  
Epoch 4/50  
32/32 [=====] - 35s 1s/step - loss: 0.5548 - acc: 0.7529 - val\_loss: 0.3994 - val\_acc: 0.8359  
Epoch 5/50  
32/32 [=====] - 33s 1s/step - loss: 0.4938 - acc: 0.7966 - val\_loss: 0.3615 - val\_acc: 0.8477  
Epoch 6/50  
32/32 [=====] - 34s 1s/step - loss: 0.4548 - acc: 0.8075 - val\_loss: 0.3842 - val\_acc: 0.8477  
Epoch 7/50  
32/32 [=====] - 34s 1s/step - loss: 0.4558 - acc: 0.8234 - val\_loss: 0.3577 - val\_acc: 0.8555  
Epoch 8/50  
32/32 [=====] - 34s 1s/step - loss: 0.3945 - acc: 0.8442 - val\_loss: 0.2646 - val\_acc: 0.9219  
Epoch 9/50  
32/32 [=====] - 35s 1s/step - loss: 0.4139 - acc: 0.8323 - val\_loss: 0.3109 - val\_acc: 0.8750  
Epoch 10/50  
32/32 [=====] - 34s 1s/step - loss: 0.3737 - acc: 0.8651 - val\_loss: 0.2602 - val\_acc: 0.9336  
Epoch 11/50  
32/32 [=====] - 33s 1s/step - loss: 0.3402 - acc: 0.8750 - val\_loss: 0.2229 - val\_acc: 0.9180  
Epoch 12/50  
32/32 [=====] - 34s 1s/step - loss: 0.3440 - acc: 0.8661 - val\_loss: 0.3537 - val\_acc: 0.8633  
Epoch 13/50  
32/32 [=====] - 33s 1s/step - loss: 0.3605 - acc: 0.8740 - val\_loss: 0.1726 - val\_acc: 0.9336  
Epoch 14/50  
32/32 [=====] - 34s 1s/step - loss: 0.3179 - acc: 0.8899 - val\_loss: 0.2144 - val\_acc: 0.9453  
Epoch 15/50  
32/32 [=====] - 33s 1s/step - loss: 0.2779 - acc: 0.8899 - val\_loss: 0.1780 - val\_acc: 0.9375  
Epoch 16/50  
32/32 [=====] - 34s 1s/step - loss: 0.2582 - acc: 0.9117 - val\_loss: 0.1756 - val\_acc: 0.9453  
Epoch 17/50  
32/32 [=====] - 33s 1s/step - loss: 0.2393 - acc: 0.9038 - val\_loss: 0.3552 - val\_acc: 0.8867  
Epoch 18/50  
32/32 [=====] - 33s 1s/step - loss: 0.2645 - acc: 0.9018 - val\_loss: 0.1999 - val\_acc: 0.9219  
Epoch 19/50  
32/32 [=====] - 34s 1s/step - loss: 0.2466 - acc: 0.9097 - val\_loss: 0.1547 - val\_acc: 0.9492  
Epoch 20/50  
32/32 [=====] - 33s 1s/step - loss: 0.2243 - acc: 0.9067 - va


l\_loss: 0.1990 - val\_acc: 0.9297  
Epoch 21/50  
32/32 [=====] - 33s 1s/step - loss: 0.2044 - acc: 0.9196 - va  
l\_loss: 0.2510 - val\_acc: 0.9023  
Epoch 22/50  
32/32 [=====] - 33s 1s/step - loss: 0.1904 - acc: 0.9296 - va  
l\_loss: 0.2234 - val\_acc: 0.9492  
Epoch 23/50  
32/32 [=====] - 33s 1s/step - loss: 0.1961 - acc: 0.9246 - va  
l\_loss: 0.7882 - val\_acc: 0.7578  
Epoch 24/50  
32/32 [=====] - 33s 1s/step - loss: 0.1636 - acc: 0.9415 - va  
l\_loss: 0.1821 - val\_acc: 0.9375  
Epoch 25/50  
32/32 [=====] - 33s 1s/step - loss: 0.1941 - acc: 0.9296 - va  
l\_loss: 0.1739 - val\_acc: 0.9453  
Epoch 26/50  
32/32 [=====] - 34s 1s/step - loss: 0.1382 - acc: 0.9464 - va  
l\_loss: 0.1607 - val\_acc: 0.9609  
Epoch 27/50  
32/32 [=====] - 34s 1s/step - loss: 0.1445 - acc: 0.9474 - va  
l\_loss: 0.2568 - val\_acc: 0.9336  
Epoch 28/50  
32/32 [=====] - 34s 1s/step - loss: 0.1949 - acc: 0.9326 - va  
l\_loss: 0.2106 - val\_acc: 0.9453  
Epoch 29/50  
32/32 [=====] - 34s 1s/step - loss: 0.1230 - acc: 0.9554 - va  
l\_loss: 0.2069 - val\_acc: 0.9375  
Epoch 30/50  
32/32 [=====] - 33s 1s/step - loss: 0.1630 - acc: 0.9415 - va  
l\_loss: 0.2407 - val\_acc: 0.9258  
Epoch 31/50  
32/32 [=====] - 35s 1s/step - loss: 0.1015 - acc: 0.9673 - va  
l\_loss: 0.2521 - val\_acc: 0.9297  
Epoch 32/50  
32/32 [=====] - 33s 1s/step - loss: 0.1205 - acc: 0.9484 - va  
l\_loss: 0.2389 - val\_acc: 0.9258  
Epoch 33/50  
32/32 [=====] - 33s 1s/step - loss: 0.1008 - acc: 0.9633 - va  
l\_loss: 0.3344 - val\_acc: 0.8477  
Epoch 34/50  
32/32 [=====] - 33s 1s/step - loss: 0.1096 - acc: 0.9613 - va  
l\_loss: 0.1751 - val\_acc: 0.9453  
Epoch 35/50  
32/32 [=====] - 33s 1s/step - loss: 0.1293 - acc: 0.9603 - va  
l\_loss: 0.1502 - val\_acc: 0.9648  
Epoch 36/50  
32/32 [=====] - 34s 1s/step - loss: 0.0819 - acc: 0.9683 - va  
l\_loss: 0.2847 - val\_acc: 0.9141  
Epoch 37/50  
32/32 [=====] - 33s 1s/step - loss: 0.0801 - acc: 0.9732 - va  
l\_loss: 0.1821 - val\_acc: 0.9531  
Epoch 38/50  
32/32 [=====] - 33s 1s/step - loss: 0.0902 - acc: 0.9633 - va  
l\_loss: 0.2095 - val\_acc: 0.9297  
Epoch 39/50  
32/32 [=====] - 33s 1s/step - loss: 0.0518 - acc: 0.9782 - va  
l\_loss: 0.2316 - val\_acc: 0.9375  
Epoch 40/50


```
32/32 [=====] - 33s 1s/step - loss: 0.0861 - acc: 0.9673 - va
l_loss: 0.3225 - val_acc: 0.9219
Epoch 41/50
32/32 [=====] - 34s 1s/step - loss: 0.0647 - acc: 0.9792 - va
l_loss: 0.2423 - val_acc: 0.9180
Epoch 42/50
32/32 [=====] - 34s 1s/step - loss: 0.0578 - acc: 0.9752 - va
l_loss: 0.2998 - val_acc: 0.9141
Epoch 43/50
32/32 [=====] - 33s 1s/step - loss: 0.0695 - acc: 0.9742 - va
l_loss: 0.6606 - val_acc: 0.7969
Epoch 44/50
32/32 [=====] - 33s 1s/step - loss: 0.0631 - acc: 0.9762 - va
l_loss: 0.3419 - val_acc: 0.9258
Epoch 45/50
32/32 [=====] - 33s 1s/step - loss: 0.0973 - acc: 0.9603 - va
l_loss: 0.2731 - val_acc: 0.9258
Epoch 46/50
32/32 [=====] - 34s 1s/step - loss: 0.0479 - acc: 0.9762 - va
l_loss: 0.2973 - val_acc: 0.9141
Epoch 47/50
32/32 [=====] - 33s 1s/step - loss: 0.0806 - acc: 0.9663 - va
l_loss: 0.2783 - val_acc: 0.9258
Epoch 48/50
32/32 [=====] - 34s 1s/step - loss: 0.0569 - acc: 0.9752 - va
l_loss: 0.2946 - val_acc: 0.9219
Epoch 49/50
32/32 [=====] - 34s 1s/step - loss: 0.0525 - acc: 0.9782 - va
l_loss: 0.2395 - val_acc: 0.9531
Epoch 50/50
32/32 [=====] - 34s 1s/step - loss: 0.0564 - acc: 0.9785 - va
l_loss: 0.2876 - val_acc: 0.9258
```

**Using Tensorflow Intel API We were able to train our model faster then usual**

In [24]:

```
1 acc = history.history['acc']
2 val_acc = history.history['val_acc']
3 loss = history.history['loss']
4 val_loss = history.history['val_loss']
5
6 epochs = range(1, len(acc) + 1)
7
8 #Train and validation accuracy
9 plt.plot(epochs, acc, 'b', label='Training accuracy')
10 plt.plot(epochs, val_acc, 'r', label='Validation accuracy')
11 plt.title('Training and Validation accuracy')
12 plt.legend()
13
14
15
16 plt.figure()
17 #Train and validation Loss
18 plt.plot(epochs, loss, 'b', label='Training loss')
19 plt.plot(epochs, val_loss, 'r', label='Validation loss')
20 plt.title('Training and Validation loss')
21 plt.legend()
22
23 plt.show()
```





---

## Validating our Initial Model

```
In [25]: 1 Y_pred = model.predict(X_val)
2 print(Y_pred.shape)
3 #y_pred = np.argmax(Y_pred, axis=1) - used for multiclass
4 y_pred = (Y_pred > 0.5) * 1.0
5 y_pred = y_pred.reshape(y_val.shape)
6 y_pred.sum()
```

```
9/9 [=====] - 3s 248ms/step
(261, 1)
```

```
Out[25]: 146.0
```

In [26]:

```
1 print('Confusion Matrix')
2 print(confusion_matrix(y_val, y_pred))
3
4
5 print('Classification Report')
6 target_names = ['Weed', 'No Weed']
7 print(classification_report(y_val, y_pred, target_names=target_names))
8
9 # accuracy: (tp + tn) / (p + n)
10 accuracy = accuracy_score(y_val, y_pred)
11 print('Accuracy: %f' % accuracy)
12 # precision tp / (tp + fp)
13 precision = precision_score(y_val, y_pred)
14 print('Precision: %f' % precision)
15 # recall: tp / (tp + fn)
16 recall = recall_score(y_val, y_pred)
17 print('Recall: %f' % recall)
18 # f1: 2 tp / (2 tp + fp + fn)
19 f1 = f1_score(y_val, y_pred)
20 print('F1 score: %f' % f1)
```

Confusion Matrix

```
[[105  17]
 [ 10 129]]
```

Classification Report

	precision	recall	f1-score	support
Weed	0.91	0.86	0.89	122
No Weed	0.88	0.93	0.91	139
accuracy			0.90	261
macro avg	0.90	0.89	0.90	261
weighted avg	0.90	0.90	0.90	261

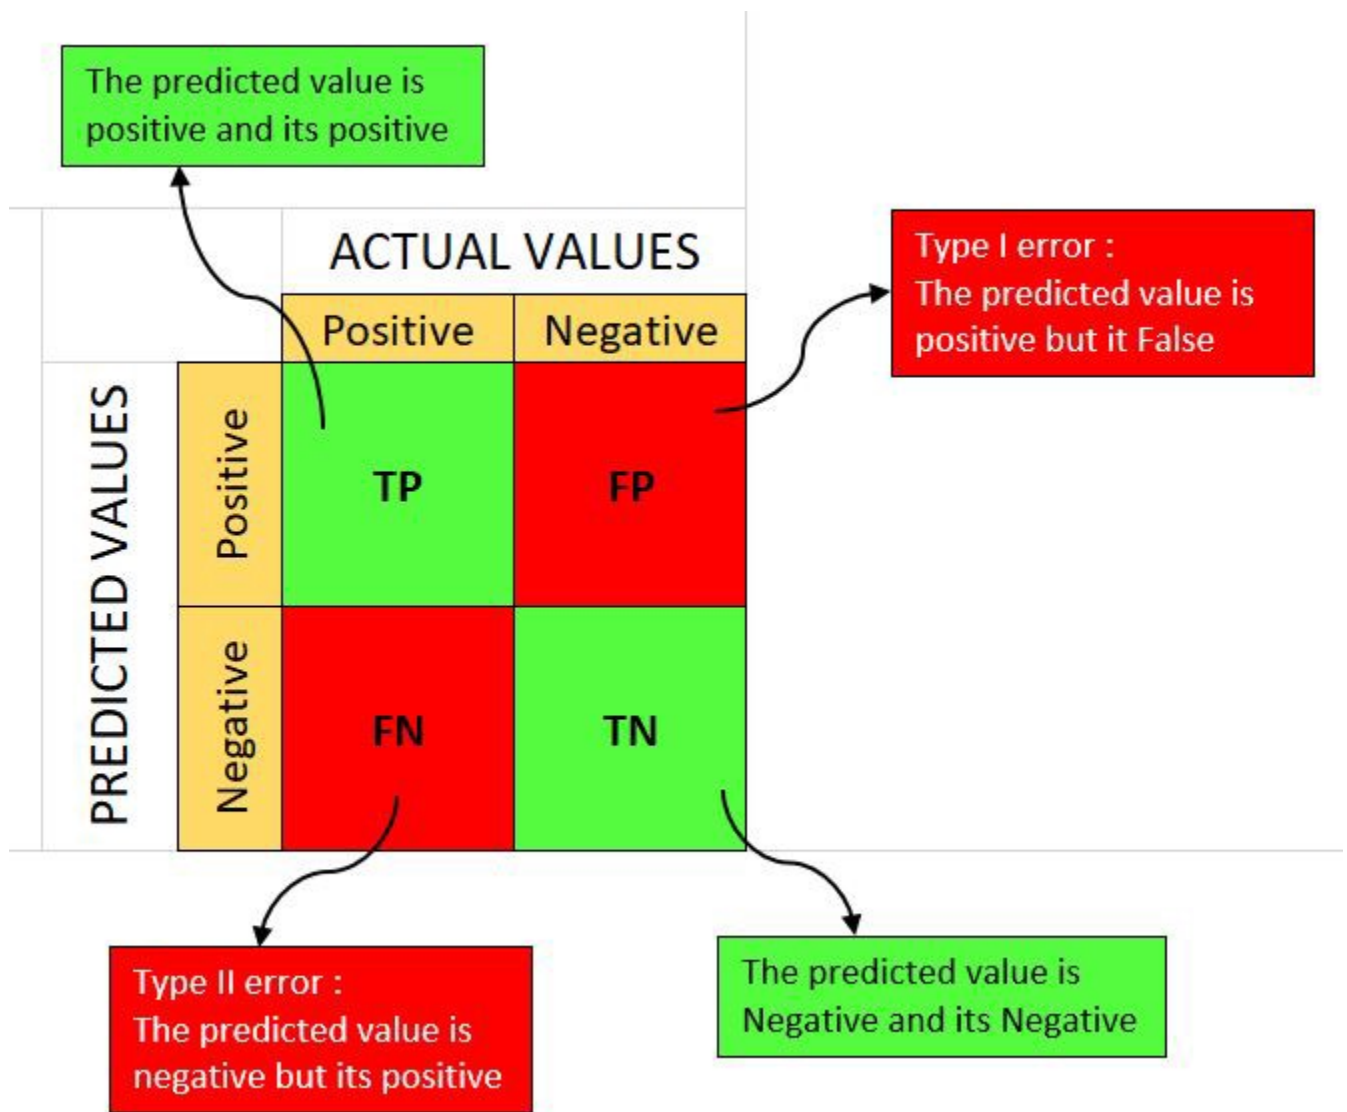
Accuracy: 0.896552

Precision: 0.883562

Recall: 0.928058

F1 score: 0.905263






---

## Transfer Learning

- Transfer learning is a machine learning technique where a model trained on one task is re-purposed on a second related task. Transfer learning can be used when the dataset is small, by using a pre-trained model on similar images we can easily achieve high performance.

# Transfer Learning



[www.analyticssteps.com](http://www.analyticssteps.com)

```
In [27]: 1 base_model = tf.keras.applications.MobileNetV2(input_shape = (224, 224, 3), include_top=False)
```

```
In [28]: 1 base_model.trainable = False
```

```
In [29]: 1 model = tf.keras.Sequential([base_model,
2                                     tf.keras.layers.GlobalAveragePooling2D(),
3                                     tf.keras.layers.Dropout(0.2),
4                                     tf.keras.layers.Dense(1, activation="sigmoid")
5                                     ])
```

---

## Training our Final Model

```
In [30]: 1 base_learning_rate = 0.00001
          2 model.compile(optimizer=tf.keras.optimizers.Adam(lr=base_learning_rate),
          3                   loss=tf.keras.losses.BinaryCrossentropy(from_logits=False),
          4                   metrics=['accuracy'])
          5
          6 history = model.fit(X_train,y_train,epochs = 50 , validation_data = (X_val, y_val))
```

WARNING:absl:`lr` is deprecated, please use `learning\_rate` instead, or use the legacy optimizer, e.g.,tf.keras.optimizers.legacy.Adam.

Epoch 1/50

33/33 [=====] - 21s 510ms/step - loss: 0.5534 - accuracy: 0.7260 - val\_loss: 0.3660 - val\_accuracy: 0.8889

Epoch 2/50

33/33 [=====] - 16s 486ms/step - loss: 0.3868 - accuracy: 0.8567 - val\_loss: 0.2973 - val\_accuracy: 0.9004

Epoch 3/50

33/33 [=====] - 16s 486ms/step - loss: 0.3325 - accuracy: 0.8798 - val\_loss: 0.2698 - val\_accuracy: 0.9004

Epoch 4/50

33/33 [=====] - 16s 497ms/step - loss: 0.3236 - accuracy: 0.8913 - val\_loss: 0.2513 - val\_accuracy: 0.9119

Epoch 5/50

33/33 [=====] - 16s 477ms/step - loss: 0.3080 - accuracy: 0.8923 - val\_loss: 0.2355 - val\_accuracy: 0.9157

Epoch 6/50

33/33 [=====] - 16s 480ms/step - loss: 0.2821 - accuracy: 0.9221 - val\_loss: 0.2111 - val\_accuracy: 0.9333

**Using Tensorflow Intel API We were able to train our model faster then usual**

In [31]:

```
1 acc = history.history['accuracy']
2 val_acc = history.history['val_accuracy']
3 loss = history.history['loss']
4 val_loss = history.history['val_loss']
5 epochs_range = range(50)
6
7 plt.figure(figsize=(15, 15))
8 plt.subplot(2, 2, 1)
9 plt.plot(epochs_range, acc, label='Training Accuracy')
10 plt.plot(epochs_range, val_acc, label='Validation Accuracy')
11 plt.legend(loc='lower right')
12 plt.title('Training and Validation Accuracy')
13
14 plt.subplot(2, 2, 2)
15 plt.plot(epochs_range, loss, label='Training Loss')
16 plt.plot(epochs_range, val_loss, label='Validation Loss')
17 plt.legend(loc='upper right')
18 plt.title('Training and Validation Loss')
19 plt.show()
```



In [32]:

```
1 Y_pred = model.predict(X_val)
2 print(Y_pred.shape)
3 #y_pred = np.argmax(Y_pred, axis=1) - used for multiclass
4 y_pred = (Y_pred > 0.5) * 1.0
5 y_pred = y_pred.reshape(y_val.shape)
6 y_pred.sum()
```

9/9 [=====] - 4s 354ms/step  
(261, 1)

Out[32]: 124.0

**Inference Time of our Final Model is 342ms/step**

## Calculating all the Evaluation Matrix

```
In [33]: 1 print('Confusion Matrix')
2 print(confusion_matrix(y_val, y_pred))
3
4
5 print('Classification Report')
6 target_names = ['No Weed', 'Weed']
7 print(classification_report(y_val, y_pred, target_names=target_names))
8
9 # accuracy: (tp + tn) / (p + n)
10 accuracy = accuracy_score(y_val, y_pred)
11 print('Accuracy: %f' % accuracy)
12 # precision tp / (tp + fp)
13 precision = precision_score(y_val, y_pred)
14 print('Precision: %f' % precision)
15 # recall: tp / (tp + fn)
16 recall = recall_score(y_val, y_pred)
17 print('Recall: %f' % recall)
18 # f1: 2 tp / (2 tp + fp + fn)
19 f1 = f1_score(y_val, y_pred)
20 print('F1 score: %f' % f1)
```

Confusion Matrix

```
[[122  0]
 [ 15 124]]
```

Classification Report

	precision	recall	f1-score	support
No Weed	0.89	1.00	0.94	122
Weed	1.00	0.89	0.94	139
accuracy			0.94	261
macro avg	0.95	0.95	0.94	261
weighted avg	0.95	0.94	0.94	261

Accuracy: 0.942529

Precision: 1.000000

Recall: 0.892086


F1 score: 0.942966

---

## Confusion Matrix

```
In [34]: 1 confusion_matrix = metrics.confusion_matrix(y_val, y_pred)
2 cm_display = metrics.ConfusionMatrixDisplay(confusion_matrix = confusion_matrix, di
```

```
In [35]: 1 cm_display.plot()  
2 plt.show()
```



---

## Weed Detection Assignment Summary

**We are able to achieve F1 Score of 0.95 with a inference time of 354ms / step with this Model**

```
In [39]: 1 f1 = f1_score(y_val, y_pred)
```

```
In [40]: 1 f1_text = f"**F1 Score:** {f1:.2f}"  
2 display(Markdown(f1_text))
```

**F1 Score: 0.94**

```
In [ ]: 1
```